



# DBSCAN and GMMs with **tidyclust** in R

Brendan Callender




Advisor: Dr. Bodwin  
Committee Members: Dr. Lund, and Dr. Dekhtyar

# Agenda


1. What is **tidyclust**?
2. What I added to **tidyclust**
  - Density-based clustering with **DBSCAN**
  - Model-based clustering with **GMMs**

# Clustering


- Unsupervised learning method
- Used to find **groupings** in data based on internal patterns



## Example Clustering Result




## Actual Penguins Species



# There are many different clustering methods!

Each has its own **strengths** and **weaknesses**



# These methods are scattered across many R packages!


R has **many ways** to do the **same thing**

- Different packages pick what works best for them
- Ex. Making a simple scatterplot

```
1 plot(y ~ x, data) # formulas
2
3 plot(data$x, data$y) # dollar signs
4
5 data %>% # tidy
5   ggplot() +
6   geom_point(aes(x = x, y = y))
```

# The **tidyclust** package

- Unifies clustering methods in R under a common interface
  - All methods use the **same syntax!**
- Like **tidymodels** but for unsupervised learning!
  - Follows conventions set in the **tidyverse**



# A quick tutorial for **tidyclust**

```
1 library(tidyclust)  
2 library(tidymodels)
```

Using k-means clustering as an example

# Creating a clustering specification

```
1 k_means_spec <- k_means(  
2   mode = "partition",  
3   engine = "dbSCAN",  
4   num_clusters = 3  
5 )
```

## Where...

- `mode` controls the behavior of the specification
- `engine` controls which underlying package implementation to use
- `num_clusters`...

# Creating a recipe

```
1 k_means_recipe <- recipe( ~ predictor1 + predictor2 + ...) |>
2   step_naomit(all_predictors()) |>
3   step_normalize(all_numeric_predictors())
```

The recipe can be used to...

1. Select which columns to use for the model using formula notation
2. Preform preprocessing steps using `step_*` () functions
  - Example. `step_pca()` and so much more!

# Creating a workflow

```
1 k_means_wflow <- workflow() |>  
2   add_model(k_means_spec) |>  
3   add_recipe(k_means_recipe)
```

The workflow is used to combine a model and a recipe together

- Allows for easy mixing and matching of different model specifications and setups

# Fitting the model

You can fit directly using the **clustering specification**

```
1 k_means_fit <- k_means_spec |>
2   fit(~predictor1 + predictor2 + ..., data)
```

Or use a **workflow**

```
1 k_means_fit <- k_means_wf |>
2   fit(data)
```

- For workflows, we just need to provide the data since the recipe controls which columns to use

# What else?

Extract key features of the model fit

```
1 k_means_fit |> extract_fit_summary()
```

Use the model to predict on new data

```
1 predict(k_means_fit, new_data)
```

Extract underlying engine object

```
1 k_means_fit |> extract_fit_engine()
```

# Model Argument Tuning

```
1 data_cvs <- vfold_cv(data, v = 5)
2
3 k_means_tune <- k_means(num_clusters = tune())
4
5 k_means_grid <- grid_regular(
6     num_clusters(c(3, 7)),
7     levels = 5
8 )
9
10 k_means_tune_res <- tune_cluster(
11     k_means_tune,
12     resamples = data_cvs,
13     grid = k_means_grid
14 )
```

This process is the  
**same** across all  
clustering  
specifications!



# What I added to **tidyclust!**

- Density-based clustering with DBSCAN
- Model-based clustering with Gaussian Mixture Models

What this took...

1. Researching each method
2. Finding a current implementation in R
3. Writing a lot of code

# Density-based Clustering with **DBSCAN**



# DBSCAN


Density-based Spatial Clustering of Applications with Noise

Arguments:


- $\epsilon$  (epsilon)
  - Controls the radius of the region used to...
    1. Compute **density estimates**
    2. Determine **connected points**
  - MinPts
    - Used as a **density threshold** to identify clusters

# Important DBSCAN Definitions

1. The  **$\epsilon$ -neighborhood** of a point is the set of all points that are within  $\epsilon$  distance from point
2. A **core point** is a point that contains at least MinPts number of points within its  $\epsilon$ -neighborhood




# The DBSCAN Algorithm



$\epsilon = 0.28$   
**MinPts = 7**


# Core Point Discovery




The fitting process begins by **scanning** through the dataset for **core points**

# Cluster Formation

Once a core point is found  
a new **cluster is formed**  
and all points within the  
 $\epsilon$ -neighborhood of the  
point are added




# Recursive Clustering Building



If a **core point** is added to the cluster, the **cluster is expanded** to all points the  $\epsilon$ -neighborhood of the core point.

# Final Cluster Assignments

This process is  
repeated until **all**  
**core points** have  
been **found**



# DBSCAN with the `dbSCAN` package

```
1 library(dbSCAN)
2 dbSCAN(x, eps, minPts = 5, weights = NULL, borderPoints = TRUE, ...)
```

Where...

- `x` is the data to perform DBSCAN on
- `eps` is the radius of the  $\epsilon$ -neighborhood used to identify core points
- `minPts` is the density threshold used to identify core points


# DBSCAN in `tidyclust`

```
1 db_clust_spec <- db_clust(  
2   mode = "partition",  
3   engine = "dbscan",  
4   radius = NULL,  
5   min_points = NULL  
6 )
```

Where...



- `radius` is the radius of the  $\epsilon$ -neighborhood used to identify core points
- `min_points` is the density threshold used to identify core points

## How `db_clust()` fits differently than `dbSCAN()`



What happens when a border point lies  
**between** core points in  
**different clusters?**



`dbSCAN()` results can differ depending on the order the data is processed




## db\_clust() Fitting Process

1. Core point discovery
2. Cluster Formation
3. Recursive Cluster building
  - Only expand clusters to other **core points**
  - **Wait** to assign **border points** until all core points have been found
4. Assign clusters to border points based on **nearest core point**


# db\_clust () Recursive Cluster Building



# db\_clust () Final Cluster Assignments



How `db_clust()`  
predicts differently than  
`dbSCAN()`



## How `db_clust()` predicts differently than `dbSCAN()`



```
1 predict(dbSCAN_fit, data, newdata)
```

For `dbSCAN()`, a new observation will be predicted to a cluster if it lies within the  **$\epsilon$ -neighborhood** of a **any point** in a cluster


```
1 predict(db_clust_fit, new_data)
```

For `db_clust()`, a new observation will be predicted to a cluster if it lies **within** the  **$\epsilon$ -neighborhood** of a **core point**

# Prediction Comparison



# Model-based Clustering with GMMs



# Gaussian Mixture Models (GMMs)

- Assumes the data is composed of **clusters** which are each **generated** from separate **multivariate Gaussian distributions**

$$f(x) = \sum_{g=1}^G p_g \Phi(x|\mu_g, \Sigma_g)$$

Where...

- $p_g$  is the weight for the gth Gaussian component
- $\mu_g$  is the mean vector for the gth Gaussian component
- $\Sigma_g$  is the variance covariance matrix for the gth Gaussian component




## Gaussian Mixture Models (cont.)

- Model-based methods can provide soft clustering labels
  - The estimated pdfs can be used to estimate the **probability** an observation belongs to each cluster



$$\hat{\gamma}_{ic} = \frac{\hat{p}_g f(x_i | \hat{\theta}_c)}{\sum_{j=1}^C \hat{p}_j f(x_i | \hat{\theta}_j)}$$

# Importance of Variance-Covariance Matrices



For Gaussian distributions,  $\Sigma$  controls the distribution...





# Shape



# Size



# Orientation




# GMM Model Specifications

- **14 different possible model specifications**
  - Circular clusters or ellipsoidal?
  - Zero or non-zero covariances?
  - Should clusters have the same shape/size/orientation?
- Commonly referenced with 3-character **model names**

Examples:  
EII, VII, EEI, EVI, EVE, EVV, VVV

# GMM Model Specifications (cont.)


EII – Circular, Equal Size



EEE – Ellipse, Equal Orientation/Shape/Size



EEI – Ellipse, Zero Covariance, Unequal Shape/Size



VVV – Ellipse, Unequal Orientation/Shape/Size



# GMMs with the `mclust` package

```
1 Mclust(data, G = 1:9, modelNames = c("EII", "VII", ..., "VVV"), ...)
```

Where...


- `data` is the data to fit a GMM to
- `G` is number of Gaussians to fit
- `modelNames` is a vector containing the model names to fit

`Mclust` will fit all combinations of `G` and `modelNames` and return best result based on BIC

# Applying tidy principles to GMM Model Names

- Make model **arguments** more **self-documenting**
  - Argument names guide user when selecting values
- Separation of **fitting** and **tuning**
  - Fit a single model with `fit()`
  - Use tuning when comparing multiple specifications

## GMM Model Specification Tree



# GMMs via `tidyclust`

```
1 gm_clust_spec <- gm_clust(num_clusters,  
2                               circular = TRUE,  
3                               zero_covariance = TRUE,  
4                               shared_orientation = TRUE,  
5                               shared_shape = TRUE,  
6                               shared_size = TRUE) %>%  
7   set_engine("mclust") %>%  
8   set_mode("partition")
```

Where...

- `circular` controls whether fitted clusters will be circular or ellipsoidal
- `zero_covariance` controls whether clusters will have zero or non-zero covariances
- `shared_{ }` controls whether clusters will have a shared shape, size, and orientation

## gm\_clust () Arguments

All arguments default to TRUE


- Arguments are named such that **TRUE** means **constraining** the model to be more **simple**
- Reduces the number of parameters than need to be estimated
- Not all datasets will be able to estimate the parameters required for the most complex model

## gm\_clust () Arguments (cont.)

- 5 TRUE/FALSE arguments
  - $2^5 = 32$  argument combinations but 14 model specifications?
- Ex. Circular clusters
  - Automatically have same shape and zero covariances!



# GMM Model Specification Tree



# Fitting with gm\_clust()


```
1 gm_clust_fit <- gm_clust_spec %>%
2   fit(~ predictor1 + predictor2 + ..., data)
```

```
1 gm_recipe <- recipe(. ~ predictor1 + predictor2 + ..., data) %>%
2   step_naomit(....)
3
4 gm_workflow <- workflow() %>%
5   add_model(gm_clust_spec) %>%
6   add_recipe(gm_recipe)
7
8 gm_clust_fit <- gm_workflow %>%
9   fit()
```

# Predicting with gm\_clust()

```
1 predict(gm_clust_fit, new_data)
```

New observations will be **predicted** to belong to the cluster in which they have the **largest probability** of belonging to





Thank you!



Especially to Dr. Bodwin!!!